

Walmart Case

```
In [ ]: # import libraries
import pandas as pd
import numpy as np
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
```

```
In [ ]: df_walmart = pd.read_csv("/Users/mojo/ML/Scaler/Projects/Walmart/walmart_data.csv")
df_walmart.head()
```

```
Out[ ]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years
0	1000001	P00069042	F	0-17	10	A	2
1	1000001	P00248942	F	0-17	10	A	2
2	1000001	P00087842	F	0-17	10	A	2
3	1000001	P00085442	F	0-17	10	A	2
4	1000002	P00285442	M	55+	16	C	4+

Missing Values, Data types and Shape of data

```
In [ ]: # Shape and size of data
df_walmart.shape
```

```
Out[ ]: (550068, 10)
```

```
In [ ]: # Missing Value
df_walmart.isna().sum()
```

```
Out[ ]: User_ID          0
Product_ID          0
Gender              0
Age                0
Occupation          0
City_Category       0
Stay_In_Current_City_Years  0
Marital_Status      0
Product_Category    0
Purchase            0
dtype: int64
```

```
In [ ]: df_walmart.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null  int64
1   Product_ID                           550068 non-null  object
2   Gender                               550068 non-null  object
3   Age                                   550068 non-null  object
4   Occupation                           550068 non-null  int64
5   City_Category                        550068 non-null  object
6   Stay_In_Current_City_Years          550068 non-null  object
7   Marital_Status                      550068 non-null  int64
8   Product_Category                    550068 non-null  int64
9   Purchase                             550068 non-null  int64
dtypes: int64(5), object(5)
memory usage: 42.0+ MB

```

- No missing values
- We can see that there are 5 numerical and 5 categorical features
- Numericals: User_ID, Occupation, Marital_Status, Product_Category and Purchase
- Categorical: Product_ID, Gender, Age(represents age group), City_Category, Stay_In_Current_City_Years
- Looking at the data we can see that the Purchase is the only column which should be in numerical.
- User_ID, Occupation, Marital_Status, Product_Category : These all should be changed into categorical data.
- We are going to change all of variables as categorical form for consistency reasons.

```

In [ ]: for col in df_walmart.columns:
        if col != 'Purchase':
            # print(col)
            df_walmart[col] = df_walmart[col].astype('category')
df_walmart.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null  category
1   Product_ID                           550068 non-null  category
2   Gender                               550068 non-null  category
3   Age                                   550068 non-null  category
4   Occupation                           550068 non-null  category
5   City_Category                        550068 non-null  category
6   Stay_In_Current_City_Years          550068 non-null  category
7   Marital_Status                      550068 non-null  category
8   Product_Category                    550068 non-null  category
9   Purchase                             550068 non-null  int64
dtypes: category(9), int64(1)
memory usage: 10.3 MB

```

Non-Graphical Analysis: Value counts and unique attributes

```
In [ ]: df_walmart.describe(include='all').T
```

	count	unique	top	freq	mean	std
User_ID	550068.0	5891.0	1001680.0	1026.0	NaN	NaN
Product_ID	550068	3631	P00265242	1880	NaN	NaN
Gender	550068	2	M	414259	NaN	NaN
Age	550068	7	26-35	219587	NaN	NaN
Occupation	550068.0	21.0	4.0	72308.0	NaN	NaN
City_Category	550068	3	B	231173	NaN	NaN
Stay_In_Current_City_Years	550068	5	1	193821	NaN	NaN
Marital_Status	550068.0	2.0	0.0	324731.0	NaN	NaN
Product_Category	550068.0	20.0	5.0	150933.0	NaN	NaN
Purchase	550068.0	NaN	NaN	NaN	9263.968713	5023.065394

- We can see the unique values of all the columns in the above table.
- We can see that in purchase mean is to the right of the median then we can say that the data is slightly right skewed.
- There are total 20 types of products.
- There are 21 types occupation available in the state.

```
In [ ]: attr = ['Gender', 'Age', 'City_Category', 'Stay_In_Current_City_Years', 'Marital_Status']
for col in attr:
    print(df_walmart[col].value_counts(normalize=True)*100)
    print("*"*50)
```

```

M      75.310507
F      24.689493
Name: Gender, dtype: float64
*****
26-35   39.919974
36-45   19.999891
18-25   18.117760
46-50    8.308246
51-55    6.999316
55+      3.909335
0-17     2.745479
Name: Age, dtype: float64
*****
B      42.026259
C      31.118880
A      26.854862
Name: City_Category, dtype: float64
*****
1       35.235825
2       18.513711
3       17.322404
4+      15.402823
0       13.525237
Name: Stay_In_Current_City_Years, dtype: float64
*****
0       59.034701
1       40.965299
Name: Marital_Status, dtype: float64
*****

```

Observations:

- Unlike the popular belief 25% of the users are female and 75% are male. Need to explore more on it
- Users in age Group 26-35 takes ~40% and kids(<18) and old people are only~ 6.75%.
- We can notice a downward trend after the age of 35.
- 35% Staying in the city from 1 year, 18% from 2 years, 17% from 3 years
- 59% of the users are single and 41 % of the users are Married.


Univariate Analysis

Outlier detection

```


In [ ]: ## Outlier detection
sns.boxplot(df_walmart['Purchase']);

```




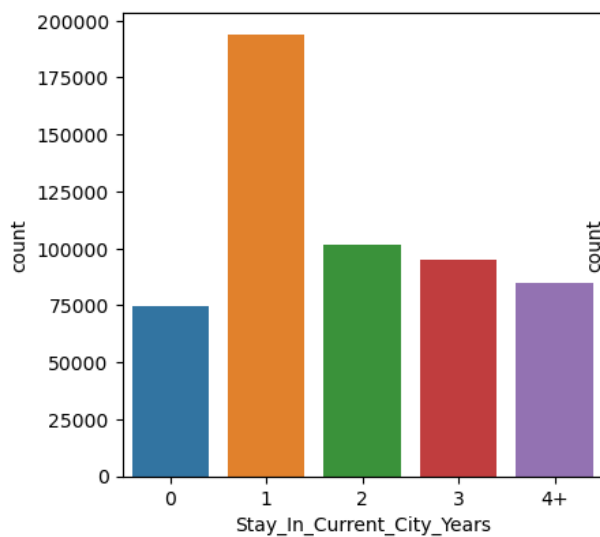




We can see that there are lots of outliers in purchase amount.


```
In [ ]: sns.histplot(df_walmart.Purchase, kde=True);
```



- Categorical variables can only be seen as count plot

```
In [ ]: fig, axs = plt.subplots(nrows=3, ncols=2, figsize=(10,15))
sns.countplot(data=df_walmart, x='Gender', ax=axs[0,0])
sns.countplot(data=df_walmart, x='Age', ax=axs[0,1])
sns.countplot(data=df_walmart, x='Occupation', ax=axs[1,0])
sns.countplot(data=df_walmart, x='City_Category', ax=axs[1,1])
sns.countplot(data=df_walmart, x='Stay_In_Current_City_Years', ax=axs[2,0])
sns.countplot(data=df_walmart, x='Marital_Status', ax=axs[2,1])
plt.show()
sns.countplot(data=df_walmart, x='Product_Category')
plt.show()
```





Observations:

- Unlike the popular belief 25% of the users are female and 75% are male. Need to explore more on it
- Users in age Group 26-35 takes ~40% and kids(<18) and old people are only~ 6.75%.
- We can notice a downward trend after the age of 35.
- 35% Staying in the city from 1 year, 18% from 2 years, 17% from 3 years
- 59% of the users are single and 41 % of the users are Married.
- Most of the users belongs to City_category B.
- Single users are more as compared to married people.
- Top 3 Product_Category: 1,5 and 11

Bivariate Analysis







We have 1- Numerical and 9 categorical Variables. For these features following plots can be drawn and analysed:


- boxplots, violinplot, and barplot
- countplot with hue

Boxplot


```
fig, axs = plt.subplots(nrows=3, ncols=2, figsize=(10,15))
count = 0
for row in range(3):
    for col in range(2):
        sns.boxplot(data=df_walmart, x=attrs[count], y='Purchase', ax= axs[r
        count+=1
plt.show()

sns.boxplot(data=df_walmart, y='Purchase', x=attrs[-1])
plt.show()
```



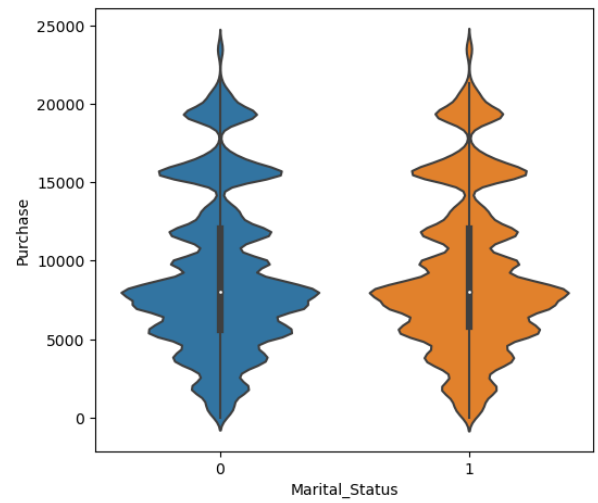
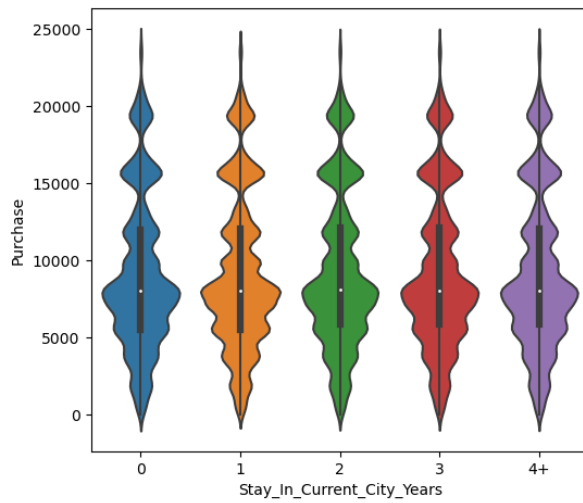
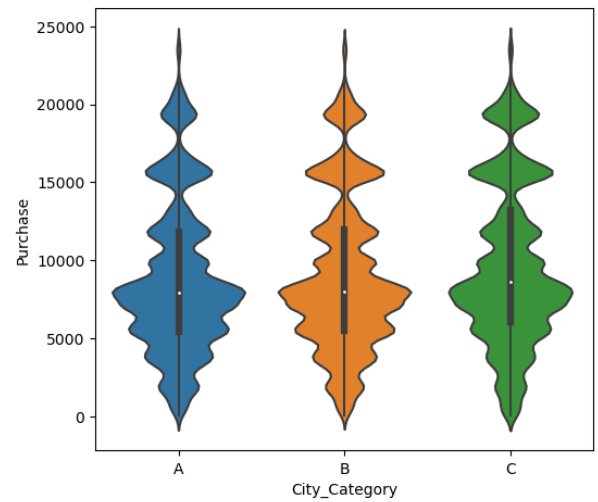






- we can see that there are outlier in all of the factors when we are comparing with respect to the Purchase amount.
- In all of the features the outliers are in the top section except the product category.
- In product category, the outliers are in bottom side of the section.

Violin plot

```
In [ ]: fig, axs = plt.subplots(nrows=3, ncols=2, figsize=(13,18))
count = 0
for row in range(3):
    for col in range(2):
        sns.violinplot(data=df_walmart, x=attrs[count], y='Purchase', ax=axs[row, col])
        count+=1
plt.show()

sns.violinplot(data=df_walmart, y='Purchase', x=attrs[-1])
plt.show()
```





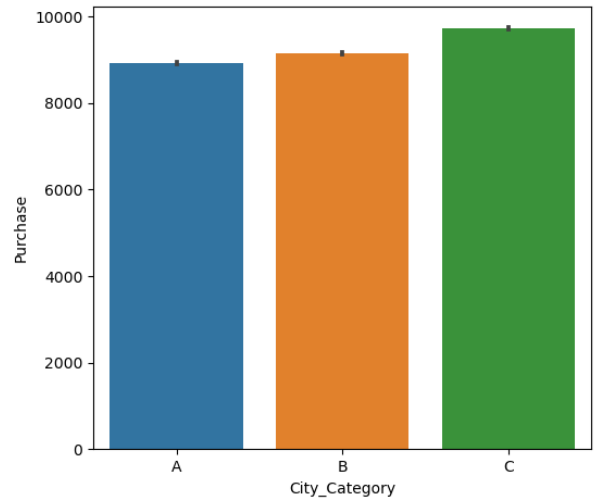






- We can notice that the distribution of all the details(e.g. male/female) of all the features are the same.
- Product category has very less distribution unable to identify as well.
- The outliers are responsible for a significant amount of purchase. This can be noticed in distribution in all the outliers and purchase.

Bar plots

```
In [ ]: fig, axs = plt.subplots(nrows=3, ncols=2, figsize=(13,18))
count = 0
for row in range(3):
    for col in range(2):
        sns.barplot(data=df_walmart, x=attrs[count], y='Purchase', ax=axs[r
        count+=1
plt.show()

sns.barplot(data=df_walmart, y='Purchase', x=attrs[-1])
plt.show()
```






- Males are purchasing more in comparison to female
- Most purchases are made among product category 10.
- Lowest purchases are in Product category: 19, 20, 12, and 13.
- Almost all of Age segments are purchasing in stores.
- Occupation 9, 19 and 20 are the ones who purchase the lowest in the stores.
- Order of Purchase City Category wise is C>B>A. This can be because of multiple factors which can be looked into more depth.
- Marital Status and # of years a person is staying in the city doesn't make any difference in this data. Which is a bit odd because if a person is married he will buy more as he lives with family.

Pair Plots

```
In [ ]: # We can see that pair plots wouldn't make any sense because we have only 1
sns.pairplot(data=df_walmart);
```



Heat map

- Heat Maps wouldn't make any sense because there is no numerical data

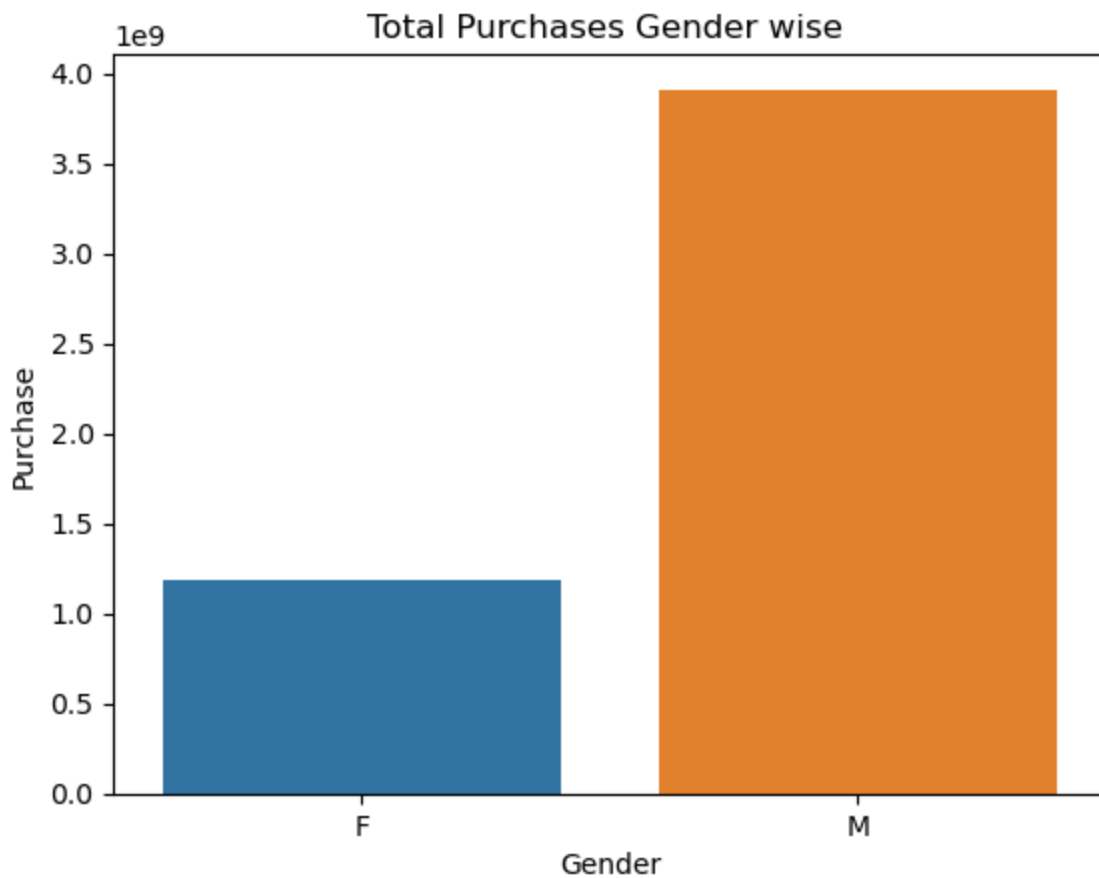
Answering questions

Q1: Are women spending more money per transaction than men? Why or Why not?

```
In [ ]: df1 = df_walmart.groupby('Gender')['Purchase'].sum().reset_index()
df1
```

```
Out[ ]:   Gender  Purchase
0      F  1186232642
1      M  3909580100
```

```
In [ ]: # sum of mal
sns.barplot(data = pd.DataFrame(df1), x='Gender', y='Purchase')
plt.title("Total Purchases Gender wise ")
plt.show()
```





```
In [ ]: df2 = df_walmart.groupby('Gender')['Purchase'].mean().reset_index()  
df2
```

```
Out[ ]:
```


	Gender	Purchase
0	F	8734.565765
1	M	9437.526040

```
In [ ]: # sum of mal  
sns.barplot(data = pd.DataFrame(df2), x='Gender', y='Purchase')  
plt.title("Average of Purchases Gender wise")  
plt.show()
```



```
In [ ]: print(df_walmart['Gender'].value_counts(normalize=True)*100)
sns.countplot(data=df_walmart, x='Gender');

M    75.310507
F    24.689493
Name: Gender, dtype: float64
```



- We can see that most of purchases are made by men in our data.
- This is because of the fact that proportions of female to male is 25%: 75% in our data.
- So we can conclude that females purchase more than males as females with 25% of population has an average of 8734.56 and males with 75% of population has average 9437.52.

Q2: Confidence intervals and distribution of the mean of the expenses by female and male customers

```
In [ ]: # number of male and female population
df_walmart['Gender'].value_counts()
```

```
Out[ ]: M    414259
        F    135809
        Name: Gender, dtype: int64
```

```
In [ ]: # df_walmart.groupby(['User_ID', 'Gender'])['Purchase'].nunique(),
df_walmart['User_ID'].nunique()
```

```
Out[ ]: 5891
```

```
In [ ]: df = pd.DataFrame(df_walmart.groupby(['User_ID', 'Gender'])['Purchase'].sum(
df_purchase_user_wise=df[df['Purchase']!=0][['Gender', 'Purchase']] # Dropping
df_purchase_user_wise
```

```
Out[ ]:
```

	Gender	Purchase
0	F	334093
3	M	810472
5	M	341635
7	M	206468
9	M	821001
...
11772	F	4116058
11774	F	1119538
11776	F	90034
11778	F	590319
11781	M	1653299


5891 rows × 2 columns

```
In [ ]: df_purchase_user_wise['Gender'].value_counts()
```

```
Out[ ]: M    4225
        F    1666
        Name: Gender, dtype: int64
```

- We can notice that we have only 5891 unique users. These consist of our male and female.


```
In [ ]: sns.kdeplot(data=df_purchase_user_wise, x='Purchase', hue='Gender', multiple
```




- We can notice that we have only 5891 unique users. These consist of our male and female.
- Female and Male purchases are right Skewed.
- Data being right skewed also points to that we have outlier in expensive purchases. This has been noticed in the above bar plots and violin plots
- We can see that there are only 4225 Males and 1666 Females.

```
In [ ]: # Creating seperate df to making counfidence interval and population mean
df_female = df_purchase_user_wise[df_purchase_user_wise['Gender']=='F']
df_male = df_purchase_user_wise[df_purchase_user_wise['Gender']=='M']
```

```
In [ ]: # Trying to find best sample size for female data size
sample_mean_trend = []
for i in np.arange(50,750):
    sample = df_female['Purchase'].sample(i)
    sample_mean_trend.append(np.mean(sample))
plt.plot(sample_mean_trend)
plt.show()
```




```
In [ ]: # Trying to find best sample size for male data size
sample_mean_trend = []
for i in np.arange(50,1000):
    sample =df_male['Purchase'].sample(i)
    sample_mean_trend.append(np.mean(sample))
plt.plot(sample_mean_trend);
```




- There is not much change after 400 for both male and female. Hence Taking 400 as sample size for male and female.

Plotting 400 samples and trying to predict population mean

```
In [ ]: num_samples = 400
female_collect_sample_means = []
for person in range(10000):
    sample_mean = df_male['Purchase'].sample(num_samples).mean()
    female_collect_sample_means.append(sample_mean)
sns.histplot(female_collect_sample_means, kde=True)
plt.title("Distribution of Means of Males with 400 Samples");
```



```
In [ ]: num_samples = 400
male_collect_sample_means = []
for person in range(10000):
    sample_mean = df_female['Purchase'].sample(num_samples).mean()
    male_collect_sample_means.append(sample_mean)
sns.histplot(male_collect_sample_means, kde=True)
plt.title("Distribution of Means of Females with 400 Samples");
```

- We can notice here that the Central Limit Theorem has been proved right. The distribution with 400 samples is gaussian in the end.

```
In [ ]: np.mean(male_collect_sample_means), np.mean(female_collect_sample_means)
print(f"Predicted Population means of the males: ${np.round(np.mean(male_col
print(f"Predicted Population means of the females: ${np.round(np.mean(female
```

```
Predicted Population means of the males: $711775.68
Predicted Population means of the females: $924637.31
```

```
In [ ]: m_upper = np.round(np.percentile(male_collect_sample_means, 97.5),2)
m_lower = np.round(np.percentile(male_collect_sample_means, 2.5),2)
f_upper = np.round(np.percentile(female_collect_sample_means, 97.5),2)
f_lower = np.round(np.percentile(female_collect_sample_means, 2.5),2)
print(f"95% Confidence Interval Population means of the males: ({m_upper}, {
print(f"95% Confidence Interval Population means of the females: ({f_upper},
```


```
95% Confidence Interval Population means of the males: (781561.77, 645098.
8)
95% Confidence Interval Population means of the females: (1018094.65, 83511
4.76)
```

Q3: Are confidence intervals of average male and female spending overlapping? How can Walmart leverage this conclusion to make changes or improvements?

```
In [ ]: fig, axs = plt.subplots()
Loading [MathJax]/extensions/Safe.js
```

```
sns.histplot(male_collect_sample_means, kde=True, legend=True)
sns.histplot(female_collect_sample_means, kde=True, legend=True)

plt.title("confidence intervals of average male and female")
plt.xlabel(['Male', 'Female']);
```



- We can clearly see that the confidence interval of Males and Females doesn't overlap.
- 95% Confidence Interval Population means of the males: (782751.49, 645250.8)
- 95% Confidence Interval Population means of the females: (1020105.31, 834682.2)
- It is very beneficial for the walmart as male and female can be easily targeted audiences are different. So walmart can easily dedicate the product in the price range as per the data. For example the females tend to spend more and they female products can be priced accordingly.


Q4: Married vs Unmarried

Are Married spending more money per transaction than unmarried? Why or Why not?

```
In [ ]: df1 = df_walmart.groupby('Marital_Status')['Purchase'].sum().reset_index()
```

```
In [ ]: # sum of mal
sns.barplot(data = pd.DataFrame(df1), x='Marital_Status', y='Purchase')
```

```
plt.title("Total Purchases Marriage wise ")
plt.show()
```




```
In [ ]: df2 = df_walmart.groupby('Marital_Status')['Purchase'].mean().reset_index()
df2
```

```
Out[ ]:
```


	Marital_Status	Purchase
0	0	9265.907619
1	1	9261.174574

```
In [ ]: sns.barplot(data = pd.DataFrame(df2), x='Marital_Status', y='Purchase')
plt.title("Average of Purchases Marital_Status wise")
plt.show()
```



```
In [ ]: print(df_walmart['Marital_Status'].value_counts(normalize=True)*100)
sns.countplot(data=df_walmart, x='Marital_Status');

0    59.034701
1    40.965299
Name: Marital_Status, dtype: float64
```



- Unmarried People has done more Purchases than married people.
- Average sales comes to be almost equal.

Confidence intervals and distribution of the mean of the expenses by Marital Status wise customers

```
In [ ]: # number of people Marital_Status wise
df_walmart['Marital_Status'].value_counts()
```

```
Out[ ]: 0    324731
        1    225337
        Name: Marital_Status, dtype: int64
```

```
In [ ]: df = pd.DataFrame(df_walmart.groupby(['User_ID', 'Marital_Status'])['Purchase']
df_purchase_marital_status_wise=df[df['Purchase']!=0][['Marital_Status', 'Purchase']]
df_purchase_marital_status_wise
```

```
Out[ ]:
```


	Marital_Status	Purchase
0	0	334093
2	0	810472
4	0	341635
7	1	206468
9	1	821001
...
11773	1	4116058
11774	0	1119538
11776	0	90034
11779	1	590319
11780	0	1653299

5891 rows × 2 columns

```
In [ ]: df_purchase_marital_status_wise['Marital_Status'].value_counts()
```

```
Out[ ]: 0    3417
        1    2474
        Name: Marital_Status, dtype: int64
```


```
In [ ]: sns.kdeplot(data=df_purchase_marital_status_wise, x='Purchase', hue='Marital_Status')
```




- There are 3417 unmarried and 2474 married people
- Data being right skewed also points to that we have outlier in expensive purchases. This has been noticed in the above bar plots and violin plots

```
In [ ]: # Creating seperate df to making counfidence interval and population mean
df_unmarried = df_purchase_marital_status_wise[df_purchase_marital_status_wi
df_married = df_purchase_marital_status_wise[df_purchase_marital_status_wise
```

```
In [ ]: # Trying to find best sample size for unmarried data size
sample_mean_trend = []
for i in np.arange(50,750):
    sample =df_unmarried['Purchase'].sample(i)
    sample_mean_trend.append(np.mean(sample))
plt.plot(sample_mean_trend)
plt.show()
```




```
In [ ]: # Trying to find best sample size for married data size
sample_mean_trend = []
for i in np.arange(50,750):
    sample =df_married['Purchase'].sample(i)
    sample_mean_trend.append(np.mean(sample))
plt.plot(sample_mean_trend)
plt.show()
```




- We can see that 300 sample size seems to be perfect

Plotting 300 samples and trying to predict population mean

```
In [ ]: num_samples = 300
unmarried_collect_sample_means = []
for person in range(10000):
    sample_mean = df_unmarried['Purchase'].sample(num_samples).mean()
    unmarried_collect_sample_means.append(sample_mean)
sns.histplot(unmarried_collect_sample_means, kde=True)
plt.title("Distribution of Means of Unmarried with 300 Samples");
```

```
In [ ]: num_samples = 300
married_collect_sample_means = []
for person in range(10000):
    sample_mean = df_married['Purchase'].sample(num_samples).mean()
    married_collect_sample_means.append(sample_mean)
sns.histplot(married_collect_sample_means, kde=True)
plt.title("Distribution of Means of married with 300 Samples");
```



```
In [ ]: np.mean(unmarried_collect_sample_means), np.mean(married_collect_sample_means)
print(f"Predicted Population means of the Unmarried: ${np.round(np.mean(unmarried_collect_sample_means), 2)}")
print(f"Predicted Population means of the Married: ${np.round(np.mean(married_collect_sample_means), 2)}")
```


Predicted Population means of the Unmarried: \$880286.01
 Predicted Population means of the Married: \$843560.42

```
In [ ]: upper1 = np.round(np.percentile(married_collect_sample_means, 97.5), 2)
lower1 = np.round(np.percentile(married_collect_sample_means, 2.5), 2)
upper2 = np.round(np.percentile(unmarried_collect_sample_means, 97.5), 2)
lower2 = np.round(np.percentile(unmarried_collect_sample_means, 2.5), 2)
print(f"95% Confidence Interval Population means of the Married: ({upper1}, {lower1})")
print(f"95% Confidence Interval Population means of the Unmarried: ({upper2}, {lower2})")
```

95% Confidence Interval Population means of the Married: (946621.04, 746166.04)
 95% Confidence Interval Population means of the Unmarried: (985301.35, 781827.96)

Are confidence intervals of average married and unmarried spending overlapping? How can Walmart leverage this conclusion to make changes or improvements?

```
In [ ]: fig, axs = plt.subplots()
sns.histplot(married_collect_sample_means, kde=True, legend=True)
sns.histplot(unmarried_collect_sample_means, kde=True, legend=True)
plt.title("confidence intervals of average married and unmarried")
plt.xlabel(['Married', "Unmarried"])
```



- Confidence interval of Married and unmarried people coincide. Which means that company can target both type of people as company is doing good to attract both of them.
- 95% Confidence Interval Population means of the Married: (946621.04, 746166.04)
- 95% Confidence Interval Population means of the Unmarried: (985301.35, 781827.96)


Q5: Results when the same activity is performed for Age

```
In [ ]: df1 = df_walmart.groupby('Age')['Purchase'].sum().reset_index()
df1
```

```
Out[ ]:
```

	Age	Purchase
0	0-17	134913183
1	18-25	913848675
2	26-35	2031770578
3	36-45	1026569884
4	46-50	420843403
5	51-55	367099644
6	55+	200767375

```
In [ ]: # sum of mal
sns.barplot(data = pd.DataFrame(df1), x='Age', y='Purchase')
plt.title("Total Purchases Age wise ")
plt.show()
```




```
In [ ]: df2 = df_walmart.groupby('Age')['Purchase'].mean().reset_index()
df2
```

```
Out[ ]:
```


	Age	Purchase
0	0-17	8933.464640
1	18-25	9169.663606
2	26-35	9252.690633
3	36-45	9331.350695
4	46-50	9208.625697
5	51-55	9534.808031
6	55+	9336.280459

```
In [ ]: # sum of mal
sns.barplot(data = pd.DataFrame(df2), x='Age', y='Purchase')
plt.title("Average Purchases Age wise")
plt.show()
```



```
In [ ]: print(df_walmart['Age'].value_counts(normalize=True)*100)
sns.countplot(data=df_walmart, x='Age');
```

```
26-35    39.919974
36-45    19.999891
18-25    18.117760
46-50     8.308246
51-55     6.999316
55+       3.909335
0-17      2.745479
Name: Age, dtype: float64
```



- Most Purchases are Made by 26-35 age group then followed by 36-45 and 18-25.

Confidence intervals and distribution of the mean of the expenses Age of the Customers

```
In [ ]: # number of cusotmers age wise
df_walmart['Age'].value_counts()
```

```
Out[ ]: 26-35    219587
        36-45    110013
        18-25     99660
        46-50     45701
        51-55     38501
        55+       21504
        0-17      15102
        Name: Age, dtype: int64
```

```
In [ ]: df_walmart['User_ID'].nunique()
```

```
Out[ ]: 5891
```

```
In [ ]: df = pd.DataFrame(df_walmart.groupby(['User_ID', 'Age'])['Purchase'].sum().r
df_purchase_age_wise=df[df['Purchase']!=0][['Age', 'Purchase']] # Dropping Use
df_purchase_age_wise
```

```
Out[ ]:
```


	Age	Purchase
0	0-17	334093
13	55+	810472
16	26-35	341635
25	46-50	206468
30	26-35	821001
...
41204	26-35	4116058
41213	46-50	1119538
41222	55+	90034
41227	46-50	590319
41232	26-35	1653299

5891 rows × 2 columns

```
In [ ]: df_purchase_age_wise['Age'].value_counts()
```

```
Out[ ]: 26-35    2053
36-45    1167
18-25    1069
46-50     531
51-55     481
55+       372
0-17      218
Name: Age, dtype: int64
```


```
In [ ]: sns.kdeplot(data=df_purchase_age_wise, x='Purchase', hue='Age', multiple='st
```



- We can see the pattern of purchase as per the age in the above graph.
- Data being right skewed also points to that we have outlier in expensive purchases. This has been noticed in the above bar plots and violin plots


```
In [ ]: # Creating seperate df to making confidence interval and population mean
df_0_17 = df_purchase_age_wise[df_purchase_age_wise['Age']=='0-17']
df_18_25 = df_purchase_age_wise[df_purchase_age_wise['Age']=='18-25']
df_26_35 = df_purchase_age_wise[df_purchase_age_wise['Age']=='26-35']
df_36_45 = df_purchase_age_wise[df_purchase_age_wise['Age']=='36-45']
df_46_50 = df_purchase_age_wise[df_purchase_age_wise['Age']=='46-50']
df_51_55 = df_purchase_age_wise[df_purchase_age_wise['Age']=='51-55']
df_55plus = df_purchase_age_wise[df_purchase_age_wise['Age']=='55+']
```

```
In [ ]: # Trying to find best sample size for age: 0_17
sample_mean_trend_0_17 = []
for i in np.arange(10,170):
    sample = df_0_17['Purchase'].sample(i)
    sample_mean_trend_0_17.append(np.mean(sample))
plt.plot(sample_mean_trend_0_17)
plt.show()
```


- Sample Size can be taken as 120

```
In [ ]: # Trying to find best sample size for age: 18-25
sample_mean_trend_18_25 = []
for i in np.arange(10,170):
    sample = df_18_25['Purchase'].sample(i)
    sample_mean_trend_18_25.append(np.mean(sample))
plt.plot(sample_mean_trend_18_25)
plt.show()
```




- Sample Size can be taken as 120

```
In [ ]: # Trying to find best sample size for age: 18-25
sample_mean_trend_26_35 = []
for i in np.arange(50,800):
    sample = df_26_35['Purchase'].sample(i)
    sample_mean_trend_26_35.append(np.mean(sample))
plt.plot(sample_mean_trend_26_35)
plt.show()
```




- Sample Size can be taken as 300

```
In [ ]: # Trying to find best sample size for age: df_36_45
sample_mean_trend_36_45 = []
for i in np.arange(50,800):
    sample = df_36_45['Purchase'].sample(i)
    sample_mean_trend_36_45.append(np.mean(sample))
plt.plot(sample_mean_trend_36_45)
plt.show()
```




- Sample Size can be taken as 450

```
In [ ]: # Trying to find best sample size for age: df_46_50
sample_mean_trend_46_50 = []
for i in np.arange(10,500):
    sample = df_46_50['Purchase'].sample(i)
    sample_mean_trend_46_50.append(np.mean(sample))
plt.plot(sample_mean_trend_46_50)
plt.show()
```




- Sample Size can be taken as 250

```
In [ ]: # Trying to find best sample size for age: df_51_55
sample_mean_trend_51_55 = []
for i in np.arange(10,450):
    sample = df_51_55['Purchase'].sample(i)
    sample_mean_trend_51_55.append(np.mean(sample))
plt.plot(sample_mean_trend_51_55)
plt.show()
```



- Sample Size can be taken as 250


```
In [ ]: # Trying to find best sample size for age: df_51_55
sample_mean_trend_55plus = []
for i in np.arange(10,300):
    sample = df_55plus['Purchase'].sample(i)
    sample_mean_trend_55plus.append(np.mean(sample))
plt.plot(sample_mean_trend_55plus)
plt.show()
```



- Sample Size can be taken as 150

Are confidence intervals of average age wise spending overlapping? How can Walmart leverage this conclusion to make changes or improvements?

```
In [ ]: fig, axs = plt.subplots()
sns.histplot(sample_mean_trend_0_17, kde=True, legend=True)
sns.histplot(sample_mean_trend_18_25, kde=True, legend=True)
sns.histplot(sample_mean_trend_26_35, kde=True, legend=True)
sns.histplot(sample_mean_trend_36_45, kde=True, legend=True)
sns.histplot(sample_mean_trend_46_50, kde=True, legend=True)
sns.histplot(sample_mean_trend_51_55, kde=True, legend=True)
sns.histplot(sample_mean_trend_55plus, kde=True, legend=True)
plt.title("confidence intervals of average age wise spending");
```



- We can see the pattern of purchase as per the age in the above graph.
- Data being right skewed also points to that we have outlier in expensive purchases. This has been noticed in the above bar plots and violin plots
- The confidence interval does collide with all the ages across. but the age group of children and older people is affected the most. Which can be a focus area for the company.

Q5: Final Insights

Observations:

- Unlike the popular belief 25% of the users are female and 75% are male. Need to explore more on it
- Users in age Group 26-35 takes ~40% and kids(<18) and old people are only~ 6.75%.
- We can notice a downward trend after the age of 35.
- 35% Staying in the city from 1 year, 18% from 2 years, 17% from 3 years
- 59% of the users are single and 41 % of the users are Married.
- Most of the users belongs to City_category B.

- Single users are more as compared to married people.
- Top 3 Product_Category: 1,5 and 11
- we can see that there are outlier in all of the factors when we are comparing with respect to the Purchase amount.
- In all of the features the outliers are in the top section except the product category.
- In product category, the outliers are in bottom side of the section.
- We can notice that the distribution of all the details(e.g. male/female) of all the features are the same.
- Product category has very less distribution unable to identify as well.
- The outliers are responsible for a significant amount of purchase. This can be noticed in distribution in all the outliers and purchase.
- We can see that in purchase mean is to the right of the median then we can say that the data is slightly right skewed.
- There are total 20 types of products.
- There are 21 types occupation available in the state.
- Males are purchasing more in comparison to female
- Most purchases are made among product category 10.
- Lowest purchases are in Product category: 19, 20, 12, and 13.
- Almost all of Age segments are purchasing in stores.
- Occupation 9, 19 and 20 are the ones who purchase the lowest in the stores.
- Order of Purchase City Category wise is $C > B > A$. This can be because of multiple factors which can be looked into more depth.
- Marital Status and # of years a person is staying in the city doesn't make any difference in this data. Which is a bit odd because if a person is married he will buy more as he lives with family.
- We can see that most of purchases are made by men in our data.
- This is because of the fact that proportions of female to male is 25%: 75% in our data.
- So we can conclude that females purchase more than males as females with 25% of population has an average of 8734.56 and males with 75% of population has average 9437.52.

- We can notice that we have only 5891 unique users. These consist of our male and female.
- Female and Male purchases are right Skewed.
- Data being right skewed also points to that we have outlier in expensive purchases. This has been noticed in the above bar plots and violin plots
- We can see that there are only 4225 Males and 1666 Females.
- We can clearly see that the confidence interval of Males and Females doesn't overlap.
- 95% Confidence Interval Population means of the males: (782751.49, 645250.8)
- 95% Confidence Interval Population means of the females: (1020105.31, 834682.2)
- It is very beneficial for the walmart as male and female can be easily targeted audiences are different. So walmart can easily dedicate the product in the price range as per the data. For example the females tend to spend more and they female products can be priced accordingly.
- Unmarried People has done more Purchases than married people.
- Average sales comes to be almost equal.
- We can see the see the pattern of purchase as per the age in the above graph.
- Data being right skewed also points to that we have outlier in expensive purchases. This has been noticed in the above bar plots and violin plots
- Teh confidence interval does collide with all the ages across. but the age group of children and older people is affected ht most. Which can be a focus area for the company.

Q6: Recommendation

- Gender wise population have differenrent habits. This can be helpful to target specific type of people.
- Males tend to spend less so they can be attracted with cheap products and females can be attracted with expensive products.
- Children, teenager and old people doesnt' spend much this can be because of the transportation reason. Because they may not able to reach to the stor.
- Company can attract them either to some kind of online store/ help them with transportation to the nearest bus stop etc.
- There are some people who are spending way more than the average people these people. Need to look into about the rest of people why is the gap between them.
- Married and unmarried people tend to spend the same so its not much useful.

- Rest of the factors don't provide much of the information. We can dig deeper with more data.

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []: